

Basics of Operating Systems

Intro:

Definition:

An OS is system software that manages computer hardware and software resources and provides common services for computer programs.

Functions:

Process Management: Manages the execution of processes, including creation, scheduling, and termination.

Memory Management: Handles allocation and deallocation of memory spaces.

File System Management: Manages files and directories, including storage, retrieval, and organization.

Device Management: Controls peripheral devices through their drivers.

Security and Access Control: Ensures authorized access to data and resources.

Types of OS:

Batch OS: Execute batches of jobs without user interaction. Suitable for repetitive tasks.

Time-Sharing OS: Allow multiple users to use the system simultaneously by quickly switching between users. Provides interactive computing.

Distributed OS: Manage a group of independent computers and make them appear to be a single coherent system.

Real-Time OS: Provide immediate processing and responses to inputs.

- **Hard Real-Time Systems:** Strict timing constraints.
- **Soft Real-Time Systems:** Less stringent timing constraints.

Network OS: Provide services to computers connected in a network, enabling resource sharing.

Mobile OS: Designed for mobile devices (e.g., iOS, Android).

Process Creation and Termination

Process Creation:

Fork System Call:

Creates a new process by duplicating an existing process.

The new process (child) is a copy of the parent process.

exec System Call:

Replaces the process's memory space with a new program.

Steps:

Assign a unique process identifier (PID).

Allocate memory for the process.

Initialize process control block (PCB).

Set up appropriate linkages (e.g., parent-child relationships).

Place the process in the ready queue.

Process Termination:

Normal Termination: Process completes its execution successfully.

Abnormal Termination: Process is terminated due to errors or external signals (e.g., divide by zero, kill signal).

Steps:

Release allocated resources.

Update process state to terminated.

Remove process from scheduling queues.

Inform the parent process (if any) of the termination.

Pre-emptive Scheduling Techniques**Definition:**

Scheduling method where the operating system can pre-empt a running process to assign the CPU to another process.

Types:

Round Robin (RR):

Each process is assigned a fixed time slice (quantum).

Processes are scheduled in a circular order, providing fair time-sharing.

Shortest Remaining Time First (SRTF):

Pre-empts the running process if a new process arrives with a shorter remaining execution time.

Aims to minimize waiting time.

Priority Scheduling:

Processes are scheduled based on priority.

The OS pre-empts the running process if a higher-priority process arrives.

Can lead to starvation of low-priority processes.

Multilevel Queue Scheduling:

Multiple queues based on priority levels, each with its own scheduling algorithm.

Pre-empts processes between different queues based on priority.

Multilevel Feedback Queue Scheduling:

Processes can move between queues based on their behaviour and requirements.

Pre-empts processes to provide flexibility and responsiveness.

Process Scheduling

Non-pre-emptive Scheduling Techniques

Definition:

Once the CPU is assigned to a process, it cannot be taken away until the process completes its execution or voluntarily releases the CPU.

Types:

First-Come, First-Served (FCFS):

Processes are scheduled in the order they arrive.

Simple but can lead to long average waiting times, particularly for processes arriving late.

Shortest Job Next (SJN):

Selects the process with the shortest execution time.

Optimal for minimizing average waiting time but requires knowledge of the process execution times.

Priority Scheduling:

Processes are scheduled based on priority.

High-priority processes are selected first.

Can lead to starvation of low-priority processes unless aging is implemented.

Race Conditions

Definition:

A situation where the behaviour and outcome of software depend on the sequence or timing of uncontrollable events, such as the order in which processes access shared resources.

Causes:

Occur when multiple processes access and manipulate shared data concurrently, leading to inconsistent or unpredictable results.

Prevention Techniques:

Mutexes:

Mutual exclusion locks to ensure that only one process can access a resource at a time.

Semaphores:

Synchronization primitives that signal and wait to control access to resources.

Monitors:

High-level synchronization constructs that encapsulate shared variables and operations, allowing only one process to execute an operation at a time.

Atomic Operations:

Operations that complete in a single step without interruption, ensuring consistency of shared data.

Threads and Multithreading

Basics of Threads:

Thread: The smallest unit of a process that can be scheduled and executed independently.

Process: Contains at least one thread.

Multithreading: Allows multiple threads to exist within a single process, sharing resources but able to run independently.

Thread Life Cycle: New -> Runnable -> Running -> Blocked -> Dead.

Advantages of Multithreading:

Resource Sharing:

Threads share the same memory space, which makes communication faster.

Responsiveness:

Applications remain responsive to user input by running separate threads for different tasks.

Utilization of Multiprocessor Architectures:

Threads can run on multiple processors, enhancing performance.

Disadvantages of Multithreading

Complexity: Writing and managing multithreaded applications can be complex.

Concurrency Issues: Problems like deadlocks, race conditions, and thread starvation can occur.

Implementing Threads

User-Level Threads: Managed by user-level libraries, fast context switching, but lack kernel support.

Kernel-Level Threads: Managed by the OS kernel, slower context switching, but better performance in multi-core processors.

Memory Management

Contiguous Memory Allocation:

Fixed Partitioning:

Memory is divided into fixed-size partitions.

Simple but can lead to internal fragmentation.

Dynamic Partitioning:

Memory is allocated dynamically, leading to external fragmentation and the need for compaction.

Paging

Paging:

Memory is divided into fixed-size pages, and physical memory is divided into frames of the same size.

Page Table:

Maintains the mapping between virtual addresses and physical addresses.

Advantages:

Eliminates external fragmentation, allows non-contiguous memory allocation.

Disadvantages:

Can introduce overhead due to the need for page table management.

Segmentation

Segmentation:

Memory is divided into segments based on logical divisions such as functions, objects, or data.

Segment Table: Maintains the base and limit addresses of each segment.

Advantages:

Provides a logical way of dividing memory, can lead to better utilization.

Disadvantages:

Can introduce external fragmentation, and requires complex segment management.

Virtual Memory

Thrashing and Page Faults:**Page Fault:**

Occurs when a requested page is not in memory, causing the OS to fetch it from the disk.

Thrashing:

A state where excessive paging operations occur, degrading performance.

Page Replacement Algorithms

First-In-First-Out (FIFO):

Replaces the oldest page. Simple but can lead to Belady's anomaly.

Least Recently Used (LRU):

Replaces the page that has not been used for the longest time. Effective but requires complex tracking.

Optimal Page Replacement:

Replaces the page that will not be used for the longest time in the future. Ideal but impractical as it requires future knowledge.

Clock Algorithm:

A circular list of pages with a reference bit to approximate LRU.

File Systems

Intro:**File System (FS):**

A method and data structure that the operating system uses to manage files on a disk or partition.

Functions of a FS:**File Creation and Deletion:**

Allowing users and applications to create and delete files.

Directory Creation and Deletion:

Organizing files into directories for easier management.

Support for Storage Devices:

Managing different types of storage devices (e.g., HDD, SSD).

File Manipulation: Providing methods to read, write, and update files.

File Protection and Security: Ensuring only authorized users can access and modify files.

File System Implementation (FSI)

- File System Structure:

Boot Block: Contains the bootstrap loader program.

Superblock: Contains information about the file system (e.g., size, number of files).

Inode Table: Contains metadata about files.

Data Blocks: Store the actual file data.

File Allocation Methods:

Contiguous Allocation:

Files are stored in contiguous blocks. Simple but can lead to fragmentation.

Linked Allocation:

Each file is a linked list of blocks. No fragmentation, but can be slow.

Indexed Allocation:

Uses an index block to keep track of all blocks belonging to a file. Balances speed and space.

Directory Implementation:

Single-Level Directory:

All files are contained in a single directory. Simple but impractical for large systems.

Two-Level Directory:

Separate directory for each user. Avoids name collisions but still limited.

Tree-Structured Directory:

Hierarchical directory structure. Most common and flexible.

Acyclic-Graph Directory:

Allows shared subdirectories and files. More complex management.

Free Space Management:

Bit Vector: Uses a bitmap to track free blocks.

Linked List: Links together all free blocks.

Grouping: Tracks free blocks in groups.

Counting: Tracks the number of contiguous free blocks.

Concurrency and Synchronization

Deadlocks

Definition:

A situation in a multitasking environment where two or more processes are unable to proceed because each is waiting for the other to release a resource.

Deadlock Prevention:

- Ensure at least one of the following conditions is impossible:

Mutual Exclusion: Not all resources are sharable.

Hold and Wait: Processes are not allowed to hold one resource while waiting for another.

No Pre-emption: Resources cannot be forcibly taken from processes.

Circular Wait: Impose an ordering on resource types and ensure processes request resources in increasing order.

Deadlock Avoidance

Banker's Algorithm:

Ensures that a system will remain in a safe state by pre-evaluating resource allocation requests.

Resource Allocation Graph:

Uses graph theory to detect the possibility of a deadlock.

Deadlock Detection and Recovery

Detection:

Wait-For Graph:

Periodically check for cycles in a graph of resource allocations.

Resource Allocation Graph:

Tracks resource allocation to detect cycles.

Recovery:

Process Termination: Abort one or more processes to break the deadlock.

Resource Pre-emption: Temporarily take resources away from processes and reassign them.

Semaphores and Mutexes**Semaphore:**

A synchronization primitive that can be used to control access to a common resource by multiple processes.

Binary Semaphore: Also called a mutex. Only two states (locked/unlocked).

Counting Semaphore: Allows more than one resource instance.

Operations:

``wait()`` (decrements the semaphore) ``signal()`` (increments the semaphore).

Mutex (Mutual Exclusion):

A locking mechanism to ensure that only one thread can access a resource at a time.

Lock/Unlock: Operations to gain and release control over the resource.

Monitors and Condition Variables

Monitor:

A high-level synchronization construct that allows safe access to shared data.

Condition Variables:

Used within monitors to manage complex synchronization.

Operations: `

wait(): (releases the monitor and waits for a condition)

signal(): (wakes up a waiting thread).

Advanced Memory Management

Cache Memory

Definition:

A small, fast memory located close to the CPU to store frequently accessed data and instructions.

Purpose: To reduce the average time to access data from the main memory.

Types of Cache:

L1 Cache: Located on the CPU chip, smallest and fastest.

L2 Cache: Larger than L1, may be on the CPU or a separate chip.

L3 Cache: Shared among cores in multi-core processors, larger but slower than L1 and L2.

Cache Replacement Policies

First-In-First-Out (FIFO):

Replaces the oldest cache line. Simple but not always efficient.

Least Recently Used (LRU):

Replaces the cache line that has not been used for the longest time.

More efficient but requires tracking of usage.

Random Replacement:

Replaces a randomly selected cache line. Simple but can be inefficient.

Least Frequently Used (LFU):

Replaces the cache line that is used least frequently.

Requires counting usage frequency.

Adaptive Replacement Cache (ARC):

Balances between LRU and LFU, adjusting dynamically to the workload.

OS Security

Access Control and Authentication

Access Control:

Mechanisms to ensure only authorized users can access resources.

Discretionary Access Control (DAC):

Access rights are assigned by the owner.

Mandatory Access Control (MAC):

Access rights are determined by the system based on policies.

Role-Based Access Control (RBAC):

Access rights are based on the roles assigned to users.

Authentication: Verifying the identity of a user or process.

Password-Based Authentication: Using passwords to authenticate users.

Multi-Factor Authentication (MFA):

Combines two or more independent credentials (e.g., password and OTP).

Biometric Authentication:

Uses biological traits (e.g., fingerprints, facial recognition).

Secure OS Design and Implementation

Principles:

Least Privilege: Users and processes operate with the minimum privileges necessary.

Défense in Depth: Multiple layers of security controls.

Fail-Safe Defaults: Access decisions should default to denial.

Security Features:

User Authentication: Ensuring only authorized users can access the system.

Access Control Mechanisms: DAC, MAC, RBAC.

Auditing and Logging: Tracking and recording system activities for monitoring and forensic analysis.

Encryption: Protecting data at rest and in transit.

Malware and Défense Mechanisms

Malware: Malicious software designed to disrupt, damage, or gain unauthorized access to systems.

Types of Malwares:

Viruses: Attach to legitimate programs and spread when executed.

Worms: Self-replicating malware that spreads across networks.

Trojan Horses: Disguised as legitimate software but perform malicious activities.

Ransomware: Encrypts data and demands payment for decryption.

Spyware: Collects information without user consent.

Adware: Displays unwanted advertisements.

Défense Mechanisms:

Antivirus Software: Detects and removes malware.

Firewalls: Monitors and controls incoming and outgoing network traffic.

Intrusion Detection Systems (IDS): Monitors network or system activities for malicious activities.

Patch Management: Regularly updating software to fix vulnerabilities.

User Education and Awareness: Training users on safe computing practices.

Intro to DBMS and RDBMS

DBMS (Database Management System)

Definition: A software system that uses a standard method of cataloguing, retrieving, and running queries on data.

Functions:

Data storage, retrieval, and update.	User and system administration.
Data security and integrity.	Backup and recovery.

Types:

Hierarchical DBMS: Data is organized in a tree-like structure.

Network DBMS: Uses a graph structure to organize data.

Relational DBMS (RDBMS): Data is stored in tables (relations).

RDBMS (Relational Database Management System)

Definition: A type of DBMS that stores data in a tabular form and allows relationships between tables through keys.

Characteristics:

Data is organized in tables.

Supports SQL for querying and managing data.

Ensures ACID properties (Atomicity, Consistency, Isolation, Durability).

Examples: MySQL, PostgreSQL, Oracle, Microsoft SQL Server.

SQL vs NoSQL

SQL (Structured Query Language)

Definition: A standard programming language for managing and manipulating relational databases.

Characteristics:

Schema-based, structured data. Relational data model.

Supports complex queries and transactions.

Examples: MySQL, PostgreSQL, SQLite.

NoSQL (Not Only SQL)

Definition: A class of database management systems that do not adhere strictly to a relational model.

Characteristics:

Schema-less, flexible data model.

Handles unstructured or semi-structured data.

High scalability and performance for large datasets.

Types:

Document Store: MongoDB, CouchDB.

Key-Value Store: Redis, DynamoDB.

Column Store: Cassandra, HBase.

Graph Store: Neo4j, ArangoDB.

SQL Queries: DDL, DML, and DCL

DDL (Data Definition Language)

Purpose: Define and modify database structure.

Commands:

CREATE: Creates a new table, view, or database.

ALTER: Modifies an existing database object.

DROP: Deletes a table, view, or database.

TRUNCATE: Removes all rows from a table without logging individual row deletions.

DML (Data Manipulation Language)

Purpose: Manipulate data within existing structures.

Commands:

SELECT: Retrieves data from a database.

INSERT: Adds new rows to a table.

UPDATE: Modifies existing data within a table.

DELETE: Removes rows from a table.

DCL (Data Control Language)

Purpose: Control access to data.

Commands:

GRANT: Gives user access privileges to the database.

REVOKE: Removes user access privileges.

SQL Queries: TCL and DQL

TCL (Transaction Control Language)

Purpose: Manage transactions in a database.

Commands:

COMMIT: Saves all changes made in the current transaction.

ROLLBACK: Undoes changes made in the current transaction.

SAVEPOINT: Sets a save point within a transaction for partial rollbacks.

DQL (Data Query Language)

Purpose: Query data from a database.

Commands:

SELECT: Retrieves data from one or more tables.

Functional Dependency, Normalization, and Normal Forms

Functional Dependency

Definition:

A relationship between two attributes, typically between a key and non-key attribute, such that the value of one attribute determines the value of the other.

Notation:

$A \rightarrow B$ (Attribute B is functionally dependent on Attribute A).

Normalization

Purpose: Organize data to reduce redundancy and improve data integrity.

Process: Decomposing a table into smaller tables and defining relationships between them.

Normal Forms

1NF (First Normal Form):

Ensures each column contains atomic (indivisible) values.

Each column contains values of a single type.

2NF (Second Normal Form):

Meets all requirements of 1NF.

All non-key attributes are fully functional dependent on the primary key.

3NF (Third Normal Form):

Meets all requirements of 2NF.

There are no transitive dependencies (non-key attribute dependent on another non-key attribute).

BCNF (Boyce-Codd Normal Form):

Meets all requirements of 3NF.

For every functional dependency $(A \rightarrow B)$, A should be a super key.